

# Databases and the Relational Data Model

Gilles Falquet

Centre universitaire d'informatique

UniGE

# Content

- Databases and Database management systems
- The Relational Data Model
- Querying Relational Databases

查询关系数据库

# Database

- A database is a collection of related data
  - known facts that can be recorded and that have implicit meaning.
- A database has the following implicit properties
  - represents some aspect of the real world (universe of discourse)
  - is a logically coherent collection of data with some inherent meaning (not random)
  - designed, built, and populated with data for a specific purpose
    - an intended group of users
    - some preconceived applications for these users

# Database management system

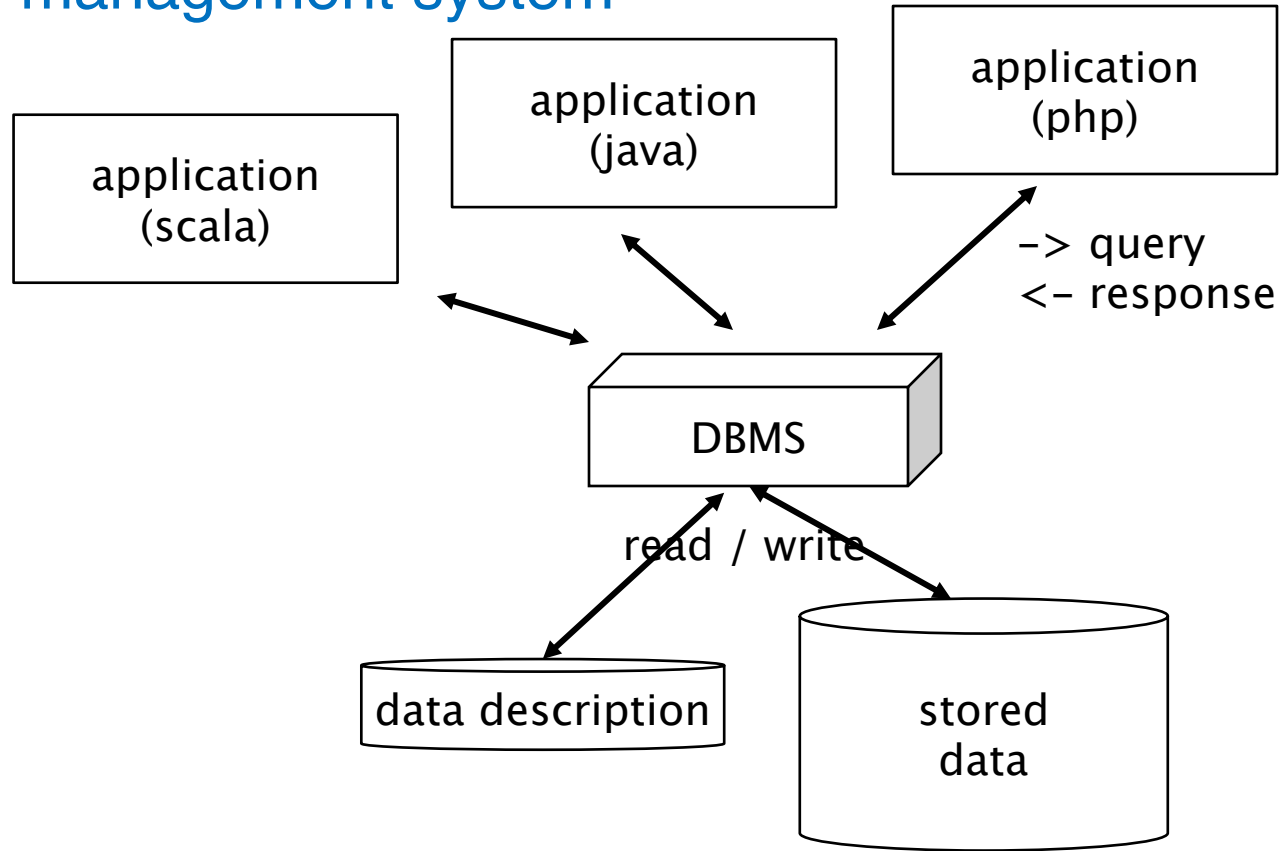
a piece of software that

- manages the physical storage of persistent data (on hard disks, solid state disks, RAM memory, etc.)
- executes application or user requests to
  - select and retrieve data
  - update data
- satisfies data management requirements

# DBMS Requirements

- persistent storage of data
- provide a description of the stored data (schema)
- data access and update functions
- content-based data retrieval (selection criteria)
- multiple simultaneous access (concurrency)
- maintain data integrity
- reliability (crash recovery mechanisms)
- prevent unauthorized access to confidential data

# Database management system



# Can't we Use a File System?

- persistent storage of data OK
- provide a description of the stored data (schema) X
- data access and update functions ~ the fs doesn't know the internal structure of a file
- content-based data retrieval (selection criteria) X
- multiple simultaneous access (concurrency) X (not within one file)
- maintain data integrity X (anything can be written in a file)
- reliability (crash recovery mechanisms) OK
- prevent unauthorized access to confidential data OK

# Data Description and Files

What is the meaning of these symbols?

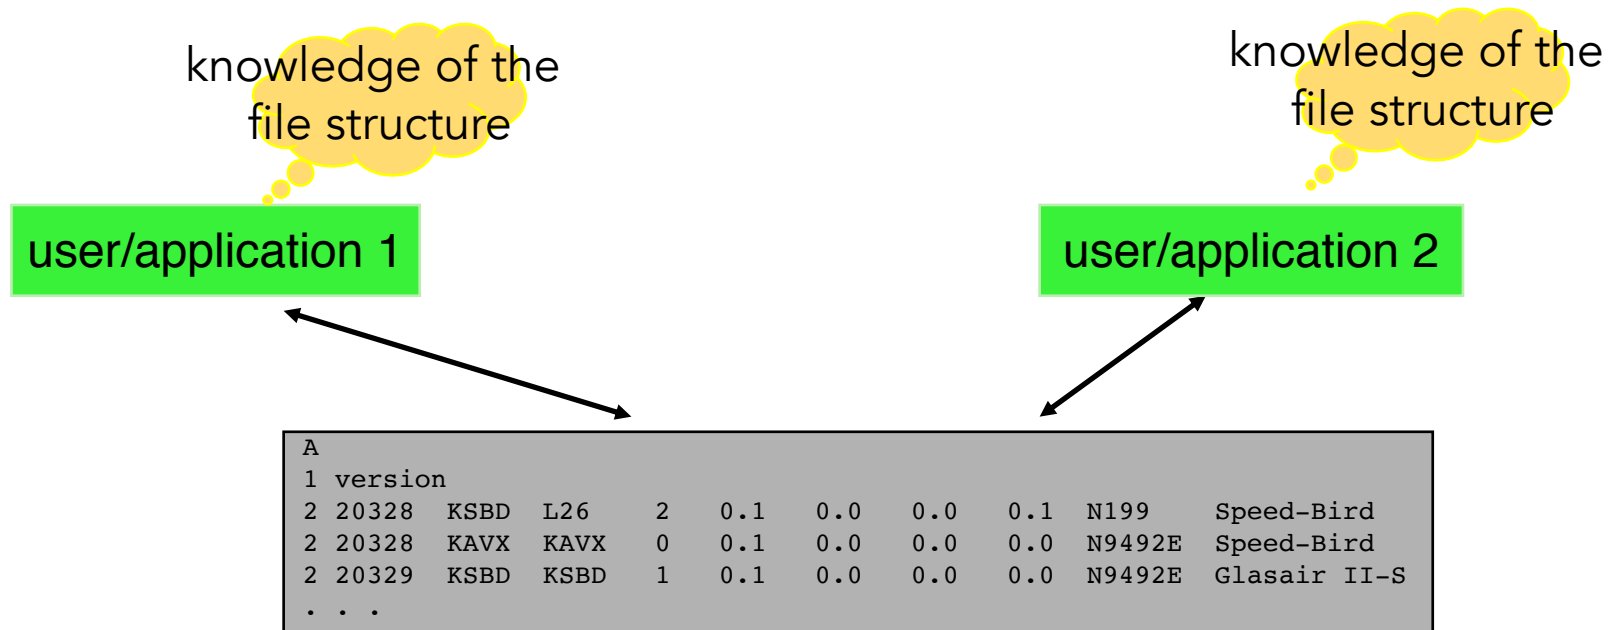
```
A
4 SoCal:hangarWall // polygon # 5
28.7 9.8 -43.7
28.7 0 -43.7
28.7 0 4
28.7 9.8 4 4
SoCal:hangarIn // polygon # 6
-28.6 9.8 -43.7
```

```
A
1 version
2 20328 KSBD L26 2 0.1 0.0 0.0 0.1 N199 Speed-Bird
2 20328 KAVX KAVX 0 0.1 0.0 0.0 0.0 N9492E Speed-Bird
2 20328 KSBD KSBD 0 0.1 0.0 0.0 0.0 N9492E Glasair II-S
2 20329 KSBD KSBD 1 0.1 0.0 0.0 0.0 N9492E Glasair II-S
. . .
```



# Problem

- Each application/user must know the meaning and organization of data in the files
- Nothing guarantees that everyone has the same interpretation



# Principle: Define a Common Schema

- A schema describes the data structure

$$\text{Database} = \text{Schema} + \text{Data}$$

- The schema is an integral part of the database
- Data may not exist without a schema

# Data


ZK567	GVA	ZRH	43	106
KL1122	AMS	CDG	50	77
KL232	AMS	PPP	560	230
LX441	GVA	NCE	35	101

## Data + Schema

Flight	From	To	Duration	Passagers
ZK567	GVA	ZRH	43	106
KL1122	AMS	CDG	50	77
KL232	AMS	PPP	560	230
LX441	GVA	NCE	35	101

# Single schema principle

- Everyone sees the data through the schema



# Data Model

- A conceptual tool to structure data
- A database schema is expressed with a data model
- A data model is based on a set of concepts
  - hierarchical models (nodes, descendant links, ...)
  - tabular models (tables, rows, columns, ...)
  - graph models (nodes, links, ...)
  - key-value (keys, value sets, ...)

# Relational Model of Data

E. F. Codd (1970)

- Based on simple mathematical constructs: sets, n-ary relations
- Has a simple intuitive interpretation: tables, rows, columns
- Well defined operations with a semantics based on relational algebra or relational calculus
- Currently the most frequently used model

# Relational Database = Set of Tables

- table = set of rows and columns
- each column has a name and a type (character string, number, date, ...)

Order

Order#	Quantity	Region	Year
1022	455	Bordeaux	2010
1322	112	Bordeaux	2010
0998	14	Valais	2014

Wine

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good



# Formal Model

**Domain (or Type):** a set of values considered as atomic

**Relation schema:**  $R(A_1, A_2, \dots, A_n)$

- a relation name  $R$
- list of attributes,  $(A_1, A_2, \dots, A_n)$ .

**Attribute:** role played by some domain  $D$  in the relation schema  $R$ .

- each  $A_k$  has a domain  $\text{dom}(A_k)$

# Relation (relation state)

for a relation schema  $R(A_1:T_1, A_2:T_2, \dots, A_n:T_n)$

a relation  $r$  on  $R$  is a set of  $n$ -tuples (the rows)

$$r = \{(t_1 = (a_1^1, \dots, a_n^1), t_2 = (a_1^2, \dots, a_n^2), \dots, t_k = (a_1^k, \dots, a_n^k))\}$$

where

$$a_1^i \in T_1, a_2^i \in T_2, \dots, a_n^i \in T_n \text{ for } i = 1, \dots, k.$$

the  $j$ th value  $a_j^i$  in tuple  $t_i = (a_1^i, \dots, a_n^i)$ , corresponds to the attribute  $A_j$ ,

notation:  $t_i.A_j$

# In SQL

- There is a set of standard types for the domains
- Definition of a relation schema (and an empty relation)

```
create table Wine(  
    Region varchar,      -- character string  
    Year integer,  
    Quality varchar,  
    AveragePrice real)
```

# Interpretation of a Relation

- Each row represents a fact about an entity or a relationship between entities
- Exemple

a row

r	y	q
---	---	---

in the table Wine(**Region**, **Year**, **Quality**) represents the fact

"Wines in region **r** have quality **q** for year **y**"

### III. Querying a Relational Database

- Extract desired information from a database
- Basic operations
  - selection, projection, join
- Standardized query language: SQL(Structured Query Language)
- A **declarative** (non-procedural language)
  - specify the result you want to obtain (what, no how)

# Selection

$r$  a relation on the schema  $(A_1, A_2, \dots, A_n)$

$$\sigma_{Condition} r = \{t \in r \mid Condition(t) = \text{true}\}$$

Retain the tuples of  $r$  that satisfy *Condition*

- the condition is a logical expression with the attribute names  $A_1, A_2, \dots, A_n$  as variables

## Exemple

$\sigma_{Quality = 'Excellent'}$ Wine

Region	Year	Quality
Bordeaux	2010	Excellent
Valais	2014	Excellent

Wine

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good

SQL:

```
select *  
from Wine  
where Quality = 'Excellent'
```

## Exemple

$\sigma_{Quality = \text{'Excellent'} \text{ and Year} > 2010}$  Wine

Région	Year	Quality
Valais	2014	Excellent

Wine

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good

```
select *  
from Wine  
where Quality = 'Excellent'  
       and Year > 2010
```



# Projection

$r$  a relation on the schema  $(A_1, A_2, \dots, A_n)$

$$B = \{A_{i_1}, A_{i_2}, \dots, A_{i_k}\} \subseteq \{A_1, A_2, \dots, A_n\}$$

$$\pi_B r = \{(t.A_{i_1}, t.A_{i_2}, \dots, t.A_{i_k} \mid t \in r \}$$

- Retain only the columns  $A_{i_1}, A_{i_2}, \dots, A_{i_k}$  of  $r$ .

# Exemple

$\pi_{\text{RegionWine}}$

Region
Bordeaux
Valais
Bourgogne

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good

## In SQL

```
select Region  
from Wine
```

Region
Bordeaux
Valais
Valais
Bordeaux
Bourgogne

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good

Duplicates are not removed !  
⇒ a multiset (bag), not a set

Reason: performance

## In SQL (2)

```
select distinct Region  
from Wine
```


Region
Bordeaux
Valais
Bourgogne

Region	Year	Quality
Bordeaux	2010	Excellent
Bourgogne	2010	Average
Valais	2014	Excellent
Valais	2013	Average
Bordeaux	2011	Good

Duplicates are removed  
(less efficient)


# Selection + Projection in SQL

```
select A1, A2, ..., Ak  
from T  
where Condition
```



## Equi Join $\bowtie$

$$r \bowtie_{A=B} s = \{t \text{ concatenated with } u \mid t \in r \text{ and } u \in s \text{ and } t.A = s.B\}$$




Student

StdNo	Name
6	Jean
3	Anne
8	Sylvie

Enrolment

Courseld	Student	Date
math	6	2019-12-04
history	3	2021-12-17
sociology	7	2019-06-06
history	6	2020-08-22



```
select * from Student, Enrolment
where Student.StdNo = Enrolment.Student
```

StdNo	Name	Courseld	Student	Date
6	Jean	math	6	2019-12-04
6	Jean	history	6	2020-08-22
3	Anne	history	3	2021-12-17

# Self Join

- A relation can be joined with itself
- Like joining two copies of the same relation

"Find the courses that share (at least) one student"


```
select E1.CourseId, E2.CourseId
from Enrolment E1, Enrolment E2
where E1.Student = E2.Student
and E1.CourseId <> E2.CourseId
```

Enrolment E1

CourseId	Student	Date
math	6	2019-12-04
history	3	2021-12-17
sociology	7	2019-06-06
history	6	2020-08-22

Enrolment E2


CourseId	Student	Date
math	6	2019-12-04
history	3	2021-12-17
sociology	7	2019-06-06
history	6	2020-08-22






## Theta Join ⋈

$$r \bowtie_{A \theta B} s = \{t \text{ concatenated with } u \mid t \in r \text{ and } u \in s \text{ and } t.A \theta s.B \}$$



# Join + Selection + Projection in SQL

```
select T1.A1, ...  
from T1, T2, ...  
where Condition and T1.Ai = T2.Bj
```



# Exercises

## Schema

- city(Name, Country, Latitude, Longitude, Population, Province)
- country(Code, Area, Capital, Name, Population, Province)
- located(City, Country, Lake, Province, River, Sea)

## Write SQL queries that retrieve the following information

- cities with more than 10 000 000 inhabitants
- countries with a city on the Indian Ocean
- cities with a population  $> 30\%$  of the country population
- name of the countries that have a city located on the Baltic Sea
- cities that are more populated than the capital of their country (\*)